

A Better Way of Organizing Your Kubernetes Manifest Files

Hugues Alary

2020/05/15

In this post I walk through the current commonly accepted way of organizing Kubernetes manifest files and propose a new approach that brings consistency, modularity and adaptability to our Kubernetes application manifests.

If you don't feel like reading a lengthy blog post and wish to get to the meat right away, [jump straight to the conclusion and my list of recommendations](#).

The current way

A common way of organizing kubernetes manifests files in a folder often looks like the following:

```
.
├── my-kubernetes-app
│   ├── config.yaml
│   ├── a-main-deployment.yaml
│   ├── a-secondary-deployment.yaml
│   ├── deployments.yaml
│   ├── services.yaml
│   ├── some-random-manifest.yaml
│   └── some-other-config.yaml
```

At first glance and for a small application this might seem like a fine structure. However, as the application grows in size and complexity, this organization quickly breaks down.

A more mature application can easily be composed of a half dozen Deployments, a few ConfigMaps, HorizontalPodAutoscalers, many Services, Ingresses, CronJobs, StatefulSets, PersistentVolumeClaims, etc.

As the number of kubernetes objects grows, so does the number of files, with each files often time grouping multiple object declarations.

For example, here's the [nginx-controller helm chart](#):

```
.
├── Chart.yaml
├── OWNERS
├── README.md
├── ci
│   └── // Folder left out for brevity
├── templates
│   ├── NOTES.txt
│   ├── _helpers.tpl
│   ├── addheaders-configmap.yaml
│   ├── admission-webhooks
│   │   └── job-patch
│   │       ├── clusterrole.yaml
│   │       ├── clusterrolebinding.yaml
│   │       ├── job-createSecret.yaml
│   │       └── job-patchWebhook.yaml
```

```

├── psp.yaml
├── role.yaml
├── rolebinding.yaml
├── serviceaccount.yaml
├── validating-webhook.yaml
├── clusterrole.yaml
├── clusterrolebinding.yaml
├── controller-configmap.yaml
├── controller-daemonset.yaml
├── controller-deployment.yaml
├── controller-hpa.yaml
├── controller-metrics-service.yaml
├── controller-poddisruptionbudget.yaml
├── controller-prometheusrules.yaml
├── controller-bsp.yaml
├── controller-role.yaml
├── controller-rolebinding.yaml
├── controller-service.yaml
├── controller-serviceaccount.yaml
├── controller-servicemonitor.yaml
├── controller-webhook-service.yaml
├── default-backend-deployment.yaml
├── default-backend-poddisruptionbudget.yaml
├── default-backend-bsp.yaml
├── default-backend-role.yaml
├── default-backend-rolebinding.yaml
├── default-backend-service.yaml
├── default-backend-serviceaccount.yaml
├── proxyheaders-configmap.yaml
├── tcp-configmap.yaml
├── udp-configmap.yaml
└── values.yaml

```

It can rapidly become hard to make sense of the overall architecture of the application: which Deployment is this Service associated with? What about this PersistentVolumeClaim, where is it being used? Does this Deployment have a HorizontalPodAutoscaler associated with it?

Often time the answer is conveyed by the filename itself. For example, a Deployment might be named `nginx-deployment.yaml` and its associated Service might be called `nginx-service.yaml`. However, with the number of files growing in the directory, it can get particularly hard to spot a file, especially when said file, actually, does not exist.

Further, the nomenclature between repositories is often times inconsistent: while, for example, some developers use the suffix `-service` to describe Service manifest files, others use `-svc`, and some prefer using a prefix.

A better way

I have been organizing my Kubernetes files in a way that I believe to be superior to the current state of affairs and my hope is that by presenting it here it might see some adoption.

Here's the basic *directory* structure:

```

├── my-kubernetes-app
│   ├── configmaps
│   ├── crons
│   ├── deployments
│   ├── hpas
│   ├── pdbs
│   ├── podpriorities
│   ├── pvcs
│   ├── services
│   ├── statefulsets
│   └── ...

```

Each line in the figure above corresponds to a directory, with each folder containing a very specific kind of Kubernetes object.

With this structure, it becomes much simpler to reason about the design of an application made of many manifest files.

Let's illustrate this with a real world example.

Below, the helm chart of the large *monolithic* application supporting Betabrand.com:

```
.
├── betabrand
│   ├── templates
│   │   ├── configmaps
│   │   │   ├── mysql.yaml
│   │   │   ├── haproxy.yaml
│   │   │   ├── magento-cron.yaml
│   │   │   ├── magento-worker.yaml
│   │   │   ├── magento.yaml
│   │   │   ├── mysql-replicator.yaml
│   │   │   ├── redis-sessions.yaml
│   │   │   └── varnish.yaml
│   │   ├── crons
│   │   │   └── database-backup.yaml
│   │   ├── deployments
│   │   │   ├── bbrecached.yaml
│   │   │   ├── haproxy.yaml
│   │   │   ├── hughbot.yaml
│   │   │   ├── imageproxy.yaml
│   │   │   ├── magento-cron.yaml
│   │   │   ├── magento-worker.yaml
│   │   │   ├── magento.yaml
│   │   │   ├── mysql-replicator.yaml
│   │   │   ├── mysql.yaml
│   │   │   ├── prerender.yaml
│   │   │   ├── pubsub-emulator.yaml
│   │   │   ├── redis-cache-primary.yaml
│   │   │   ├── redis-cache-replica.yaml
│   │   │   ├── redis-cache-sentinel.yaml
│   │   │   ├── redis-sessions-primary.yaml
│   │   │   ├── redis-sessions-replica.yaml
│   │   │   ├── redis-sessions-sentinel.yaml
│   │   │   └── varnish.yaml
│   │   ├── hpas
│   │   │   ├── magento-workers.yaml
│   │   │   └── magento.yaml
│   │   ├── pvcs
│   │   │   ├── local-dev.yaml
│   │   │   ├── mysql-replicator.yaml
│   │   │   ├── mysql.yaml
│   │   │   ├── redis-sessions-primary.yaml
│   │   │   └── redis-sessions-replica.yaml
│   │   ├── secrets
│   │   │   ├── cloudflare-api.yaml
│   │   │   ├── docker-registry.yaml
│   │   │   ├── magento-cloud-sql-service-account.yaml
│   │   │   ├── magento-global-service-account.yaml
│   │   │   ├── mysql-replicator.yaml
│   │   │   └── ssl-certificate.yaml
│   │   └── services
│   │       ├── gluster.yaml
│   │       ├── haproxy.yaml
│   │       ├── hughbot.yaml
│   │       └── imageproxy.yaml
```

```

├── magento.yaml
├── mysql-replicator.yaml
├── mysql.yaml
├── prerender.yaml
├── pubsub-emulator.yaml
├── redis-cache-primary.yaml
├── redis-cache-replica.yaml
├── redis-cache-sentinel.yaml
├── redis-cache.yaml
├── redis-sessions-primary.yaml
├── redis-sessions-replica.yaml
├── redis-sessions-sentinel.yaml
├── redis-sessions.yaml
├── varnish.yaml
├── statefulsets
│   └── gluster.yaml
├── NOTES.txt
├── _helpers.tpl
├── Chart.yaml
├── values.yaml
└── README.asciidoc

```

In the figure above, the overall architecture of the application is distinguishable at a glance as well as the software stack: mysql, haproxy, magento, bbrecached (an internal software), hughbot (another internal tool), imageproxy, prerender, google pub-sub emulator, redis, varnish and gluster.

It is made obvious that to the 18 Deployments correspond 18 Services, 8 ConfigMaps are easily matched to their 8 Deployments, a lone StatefulSet does not have any corresponding configuration, etc.

There are 2 HorizontalPodAutoscalers; can you guess which Deployments they manage autoscaling for?

Also, omitted from the previous figure but actually present in the real Betabrand repository: a `cluster-config` directory describing Kubernetes objects that I consider aren't part of the application itself as they are concerned about the cluster configuration itself. The same structure applies here too.

```

├── cluster-config
│   ├── clusterrolebindings
│   │   ├── external-dns.yaml
│   │   └── tiller.yaml
│   ├── clusterroles
│   │   └── external-dns.yaml
│   ├── deployments
│   │   └── external-dns.yaml
│   ├── podpriorities
│   │   └── high-priority.yaml
│   ├── serviceaccounts
│   │   └── external-dns.yaml
│   ├── storageclasses
│   │   ├── persistent-fast.yaml
│   │   └── persistent.yaml
│   └── README.asciidoc

```

Conclusion

So here it is, my recommendations when it comes to organizing your Kubernetes manifests:

- **Group** manifest files in **directories named after the *Kind of object***: `deployments`, `configmaps`, `services`, etc. Note the directory name lower cased and pluralized.

```

├── my-kubernetes-app

```

- configmaps
- crons
- deployments
- hpas
- pdbs
- podpriorities
- pvcs
- services
- statefulsets
- ...

- **Shorten** very long names with **commonly used abbreviations**, for example, `hpas/` instead of `horizontalpodautoscalers/`, `pdbs/` instead of `poddissruptionbudgets/`. Following the `kubectl` nomenclature is probably a good idea.
- Use **consistent** manifest filenames between `kinds/` directories. For example, if your app is named `funnygifs-slackbot`:
 - the Deployment should be named `funnygifs-slackbot.yaml`
 - the Service `funnygifs-slackbot.yaml`
 - the ConfigMap `funnygifs-slackbot.yaml`
 - the HPA `funnygifs-slackbot.yaml`
 - the PVC `funnygifs-slackbot.yaml`
 - etc.
- **Avoid stutter**: do not call a deployment manifest `funnygifs-slackbot-deploy.yaml` or a service `funnygif-slackbot-service.yaml`, etc.
- Use **resource names consistent with the manifest filename**:

```
deployments/funnygifs-slackbot.yaml
```

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: funnygifs-slackbot (1)
  labels:
    app: funnygifs-slackbot
```

1. `name:` matches the filename

```
services/funnygifs-slackbot.yaml
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: funnygifs-slackbot (1)
```

1. `name:` matches the filename

```
configmaps/funnygifs-slackbot.yaml
```

```
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: funnygifs-slackbot (1)
```

1. `name:` matches the filename

- etc.

