

# Directory structure for project with Dockerfile, Jekinsfile, Kubernetes deployment yaml, pip requirements.txt, and test scripts?

Asked 4 years, 11 months ago    Modified 4 years, 10 months ago    Viewed 3k times

▲

2

▼

★

2

🕒

Would the following directory structure work?

The goal is to have Jenkins trigger off GitHub commits and run Multi-branch Pipelines that build and test containers. (I have everything running on Kubernetes, including Jenkins)

```
/project
.git
README.md
Jenkinsfile
/docker_image_1
  Dockerfile
  app1.py
  requirements.txt
/unit_tests
  unit_test1.py
  unit_test2.py
/docker_image_2
  Dockerfile
  app2.py
  requirements.txt
/unit_tests
  unit_test1.py
  unit_test2.py
/k8s
  /dev
    deployment.yaml
  /production
    deployment.yaml
/component_tests
  component_tests.py
```


- 1. Is the k8s folder that has the deployment.yamls in the right place?
- 2. Are the test folders in good locations? The tests in "component\_tests" will ideally be doing more end-to-end integrated testing that involve multiple containers
- 3. I see a lot of repos have Jenkins file and Dockerfile in the same directory level. What are the pros and cons of that?

[docker](#) [jenkins](#) [pip](#) [kubernetes](#) [jenkins-pipeline](#)

ShareImprove this questionFollow

edited Jul 26, 2017 at 18:12

asked Jul 26, 2017 at 17:13

 [gunit](#)

3,262 ● 2 ● 26 ● 41

2 Answers

Sorted by:

Trending sort available ⓘ

Highest score (default) ⚡

▲

5

▼

🕒

**Your privacy**

By clicking "Accept all cookies", you agree Stack Exchange can store cookies on your device and disclose information in accordance with our [Cookie Policy](#).

There's no good answer to this question currently.

[Accept all cookies](#) [Customize settings](#)

Kubernetes is a container management system, but as a technology it relies on additional 3rd party tooling manage the build part of the ALM workflow. There a lots of options available for turning your source code into a container running on Kubernetes. Each has it's own

consequences for how your source code is organised and how a deployment might be invoked from a CI/CD server like Jenkins.

I provide the following collection of options for your consideration, roughly categorized. Represents my current evaluation list.

"Platform as a service" tools

Tooling the manages the entire ALM lifecycle of your code. Powerful but more complex and opinionated.

- [Deis workflow](#)
- [Openshift](#)
- [Fabric8 \(See also Openshift.io\)](#)

Build and deploy tools

Tools useful for the code/test/code/retest workflow common during development. Can also be invoked from Jenkins to abstract your build process.

- [Draft](#)
- [Forge](#)
- [Kcompose](#)
- [Fabric8 Maven plugin \(Java\)](#)
- [Psykube](#)

YAML templating tools

The kubernetes YAML was never designed to be used by human beings. Several initiatives to make this process simpler and more standardized.

- [Helm](#)
- [Ksonnet](#)

Deployment monitoring tools

These tools have conventions where they expect to find Kubernetes manifest files (or helm charts) located in your source code repository.

- [Keel](#)
- [Kube-applier](#)
- [Kubediff](#)
- [Landscaper](#)
- [Kit](#)

CI/CD tools with k8s support

- [Spinnaker](#)
- [Gitlab](#)
- [Jenkins + Kubernetes CI plugin](#)
- [Jenkins + Kubernetes plugin](#)

ShareImprove this answerFollow

edited Sep 7, 2017 at 9:03

answered Jul 26, 2017 at 18:24



Mark O'Connor

74.6k ● 10 ● 133 ● 178

▲

4

▼

✓

🕒

This is really left much to your preference. In our projects we tend to split services into separate repositories not subfolders, but we also had a

This is really left much to your preference. In our projects we tend to split services into separate repositories not subfolders, but we also had a case where we had a bunch of Scala microservices managed in similar way (although dockers were built with sbt plugin for docker)

One big advice I would give you is that in the long run managing your kubernetes manifests like that might become serious pain in the back. I went through this, and my suggestion is to use `helm` charts from day one.

I assume that your "component\_tests" are end-to-end tests. Other than naming I see no problem with that. For cases where we test solutions that span multiple repos we keep them in a separate repo as well though.

ShareImprove this answerFollow

answered Jul 26, 2017 at 17:59



Radek 'Goblin'  
Pieczonka

19.4k 5 46 45

- Definitely split services into separate repositories when possible. Then you put the deployment stuff in a separate repos as well that uses the built versions of the images from your image repository. It's so much cleaner. – [Grimmy](#) Jul 26, 2017 at 18:14
- ... at least when you work with more than a couple of people. I also don't want to see commits about configuration files for live mixed in with the apps. – [Grimmy](#) Jul 26, 2017 at 18:20
- Splitting services up in the separate repositories is a good idea and inline with the concept of microservices and the 12 Factor app recommendations. This is where Helm is a good candidate for future success with the emergence of repositories for hosting helm charts. Should make the deployment of complex eco-systems of micro services more palatable – [Mark O'Connor](#) Jul 26, 2017 at 18:37